

# Lab1 deliverable

Daniel Sattler and Cristina Raluca Vijulie

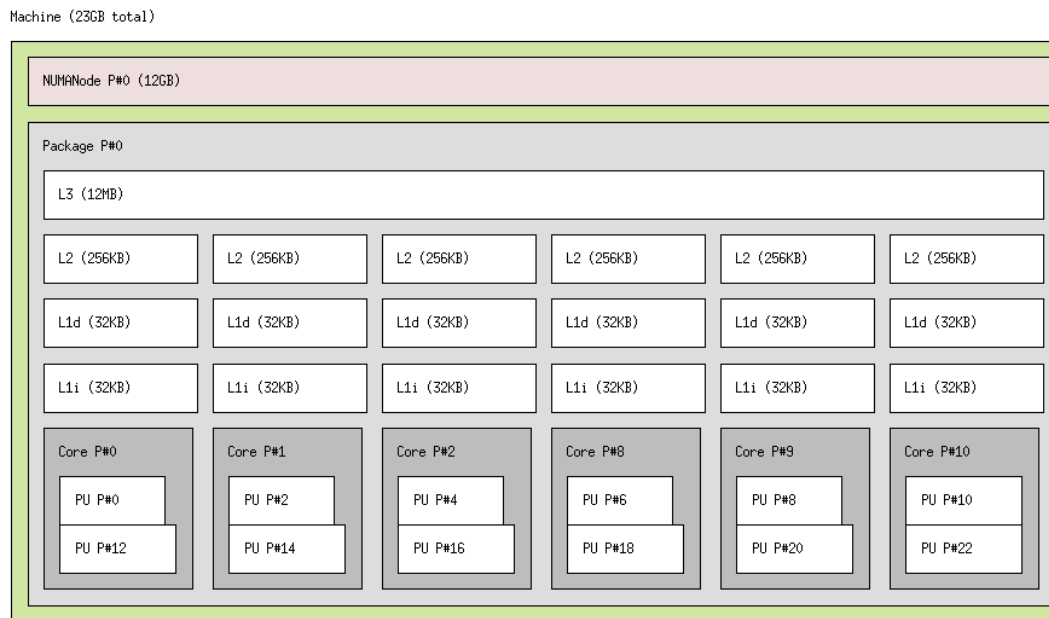
par1207

October 13, 2017

## Node architecture and memory

1. Draw and briefly describe the architecture of the computer in which you are doing this lab session (number of sockets, cores per socket, threads per core, cache hierarchy size and sharing, and amount of main memory).

- (a) sockets: 2
- (b) cores/socket: 6
- (c) threads: 2



## Timing sequential and parallel executions

2. Describe what do you need to add to your program to measure the elapsed execution time between a pair of points in the program, clearly indicating the library header file

that needs to be included, the library functions that need to be invoked, the data structure and its fields.

To measure the time we need to get a timestamp with the time of the day when the execution starts and subtract that value to the time of day when the program ends. To do so we use the following:

- The library:

```
#include <sys/time.h>
```

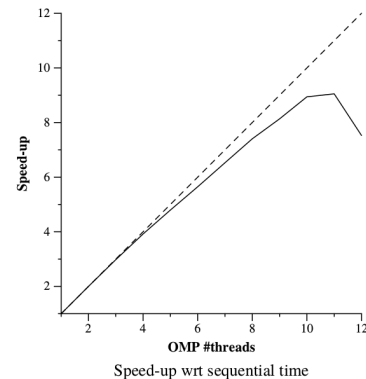

- The function:

```
int gettimeofday(struct timeval *, void *);
```

- The data structure **timeval**, which includes the following fields:

time_t	tv_sec	seconds
suseconds_t	tv_usec	microseconds

3. Plot the speedup obtained when varying the number of threads (strong scalability) and problemsize (weak scalability) for pi omp.c. Reason about how the scalability of the program.



When we are in weak scalability, the speedup does not change so much, it means the scalability is good and we can keep increasing the problem size while increasing the threads and the execution time will not be much greater. In the case of strong scalability, we can see that more than 11 threads start to decrease the speed-up, so it does not have a great strong scalability, with 11 as the optimal number of threads.


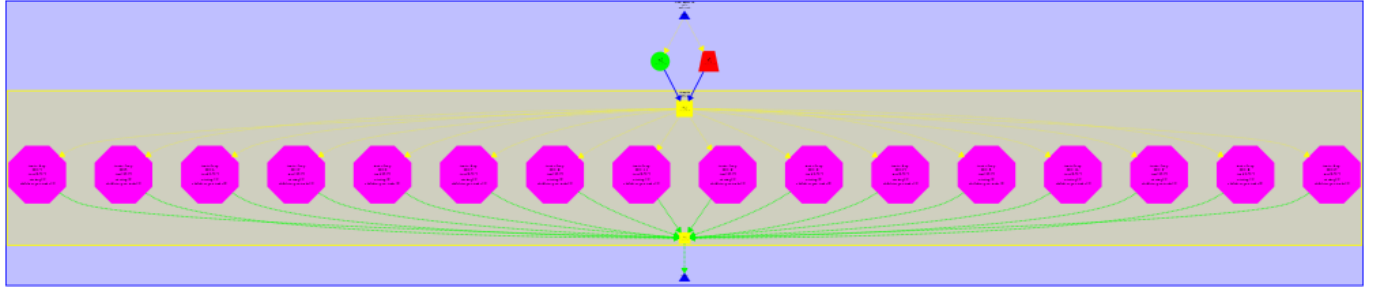
## Visualizing the task graph and data dependences

4. Include the source code for function dot product in which you show the Tareador instrumentation that has been added to study the potential parallelism in the code. This

instrumentation has to appropriately define tasks and filter the analysis of variable(s) that cause the dependence(s).

The source code has been added to the archive and is named `dot_product.c`

5. Capture the task dependence graph for that task decomposition and the execution timelines (for 8 processors) that allow you to understand the potential parallelism attainable. Briefly comment the relevant information that is reported by the tools.



## Analysis of task decompositions

6. Complete the following table for the initial and different versions generated for `3dfft_seq.c`, briefly commenting the evolution of the metrics with the different versions.

Version	$T_1$	$T_\infty$	Parallelism
seq	593772	593758	1.000
v1	593772	593758	1.000
v2	594054	315537	1.883
v3	594054	108701	5.465
v4	594054	59694	9.952

7. With the results from the parallel simulation with 2, 4, 8, 16 and 32 processors, draw the execution time and speedup plots for version v4 with respect to the sequential execution (that you can estimate from the simulation of the initial task decomposition that we provided in 3dfft seq.c, using just 1 processor). Briefly comment the scalability behaviour shown on these two plots.




Figure 1: Execution time in ns




Figure 2: Speedup v4 vs. seq

As we can see in the plots, the scalability is bad. The maximum number of cores we can use at the same time are 16, more than that do not help the execution time.

## Tracing sequential and parallel executions

8. From the instrumented version of `pi seq.c`, and using the appropriate Paraver configuration file, obtain the value of the parallel fraction for this program when executed with 100.000.000 iterations, showing the steps you followed to obtain it. Clearly indicate which Paraver configuration file(s) did you use.




Figure 3: Parallel fraction of the code

The parallel fraction of this program is the parallel duration \* 100 / the total duration. The value of the fraction is: 58.4% This was obtained using the `APP_userevents` configuration.

9. From the instrumented version of `pi omp.c`, and using the appropriate Paraver configuration file, show a profile of the % of time spent in the different OpenMP states when using 8 threads and for 100.000.000 iterations. Clearly indicate which Paraver configuration file(s) did you use and your own conclusions from that profile.

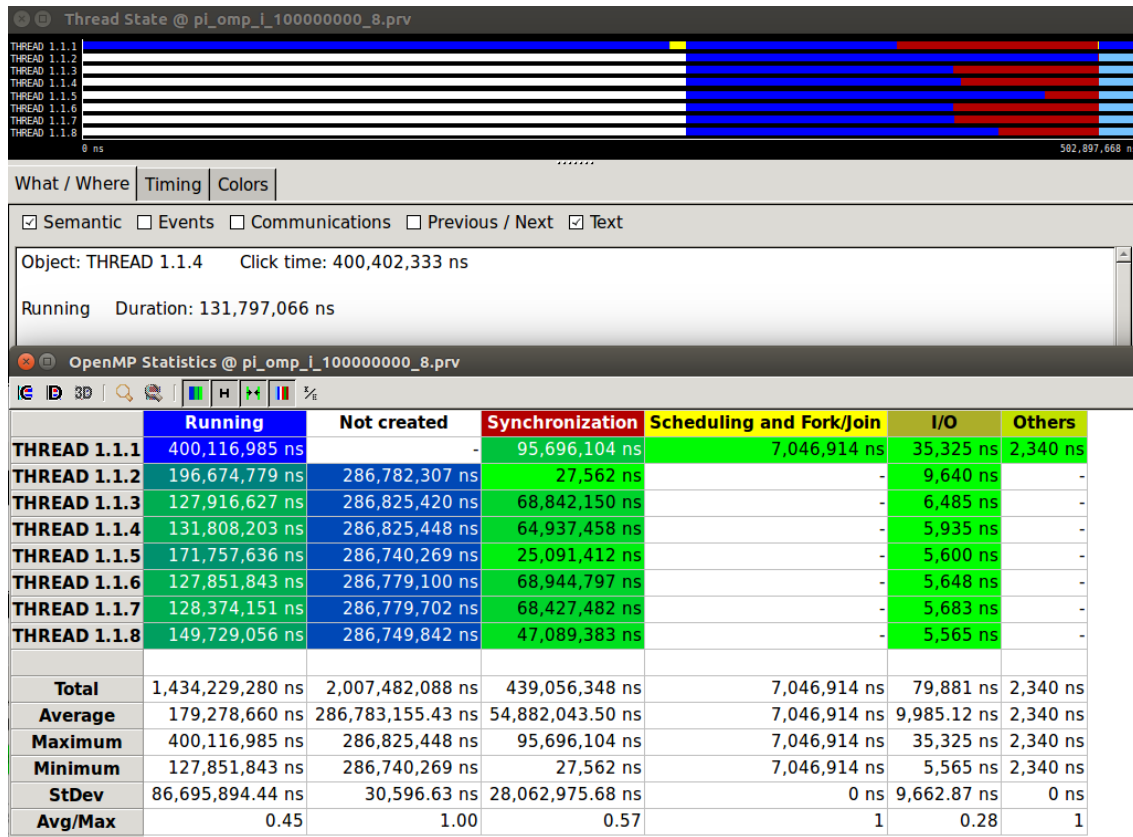


Figure 4: Time spent in the different states of OMP

This was obtained using the OMP\_state\_profile configuration.